Deep Reinforcement Learning
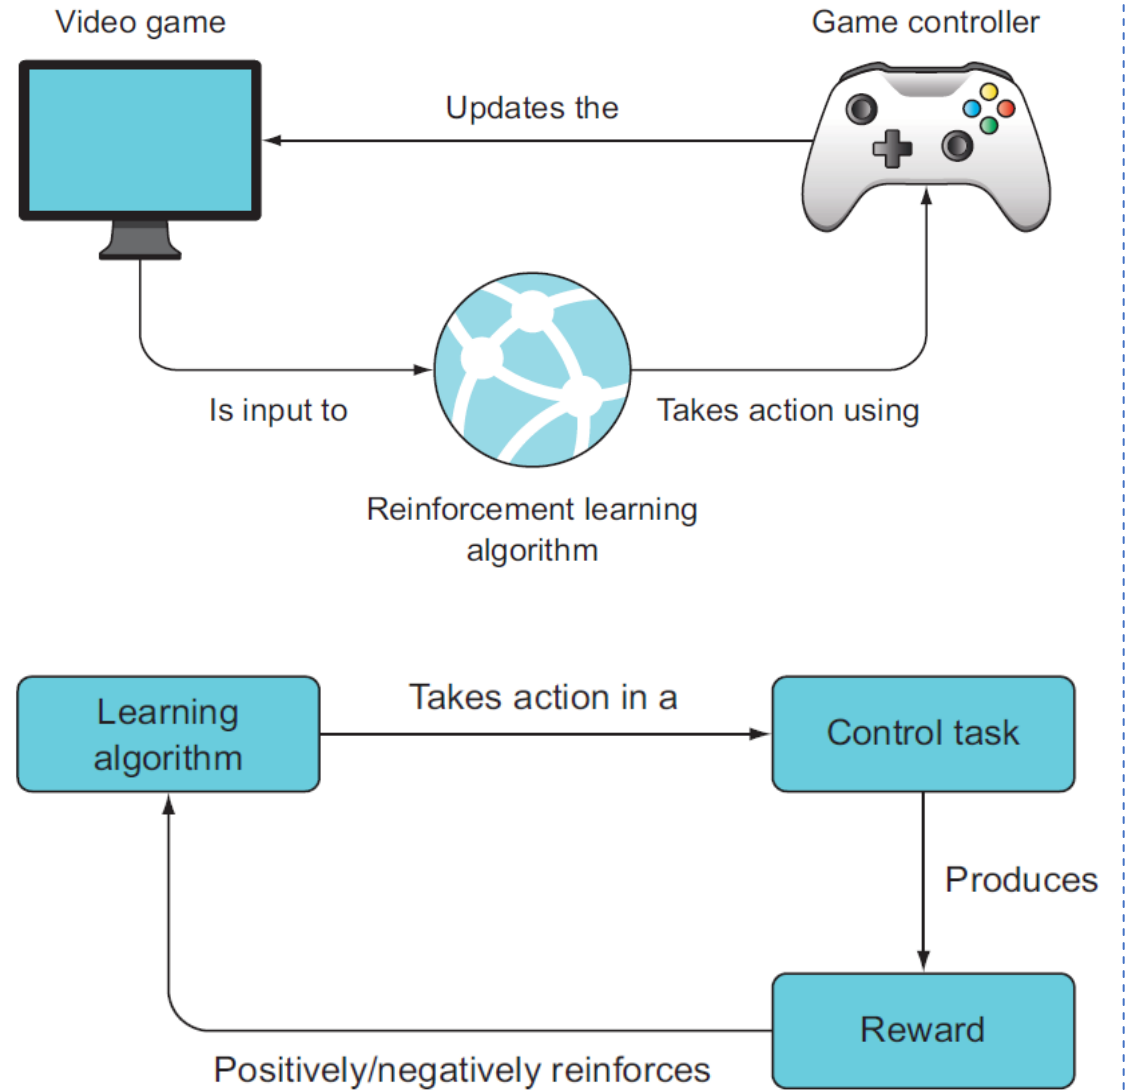
Kalle.Prorok@gmail.com , Umeå 2020

Parts based on the book Deep Reinforcement Learning in Action By Alexander Zai and Brandon Brown

And the free Sutton-Barto book Reinforcement Learning: An Introduction

Reinforcement learning

Deep Reinforcement Learning: Pong from Pixels

REINFORCEMENT LEARNING
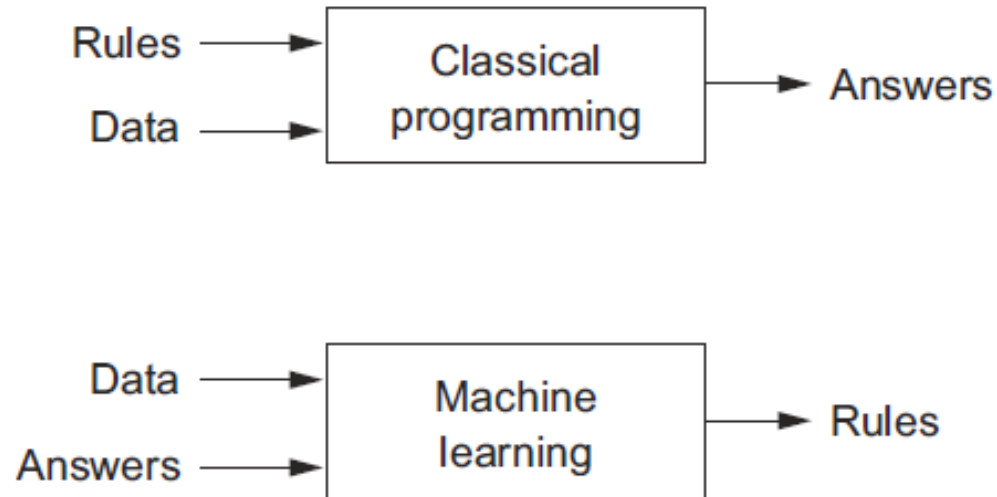

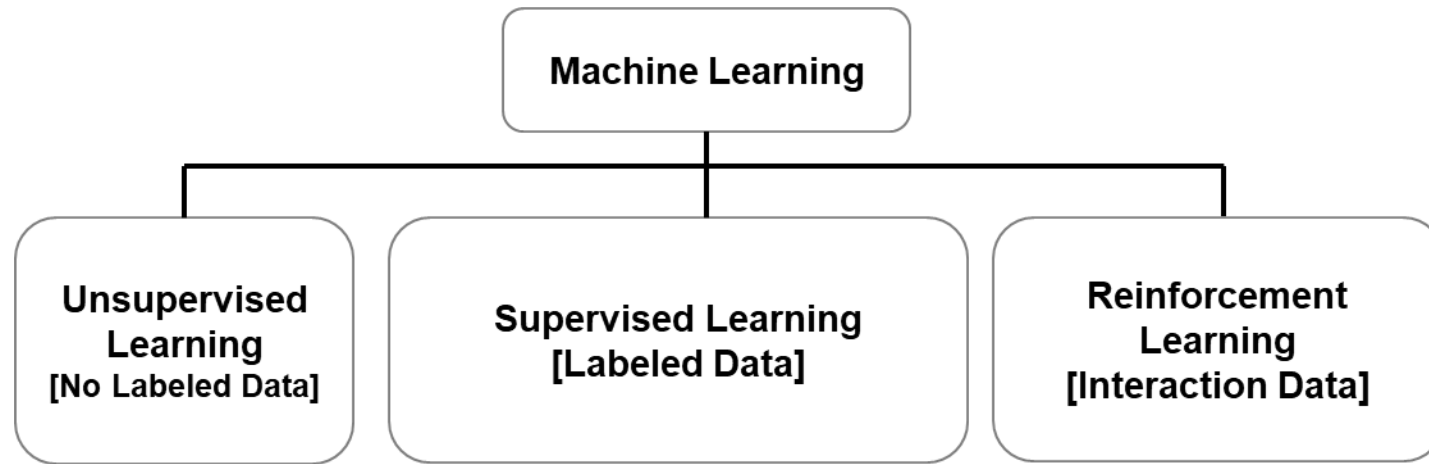Ex. Smart Elevators

NEW WAY TO PROGRAM

Figure 1.2 Machine learning: a new programming paradigm

Ur F Chollet: Deep Learning with Python

UMEÅ UNIVERSITET

# Reinforcement Learning vs Machine Learning vs Deep Learning

Machine Learning

Unsupervised Learning
[No Labeled Data]

Supervised Learning
[Labeled Data]

Reinforcement Learning
[Interaction Data]
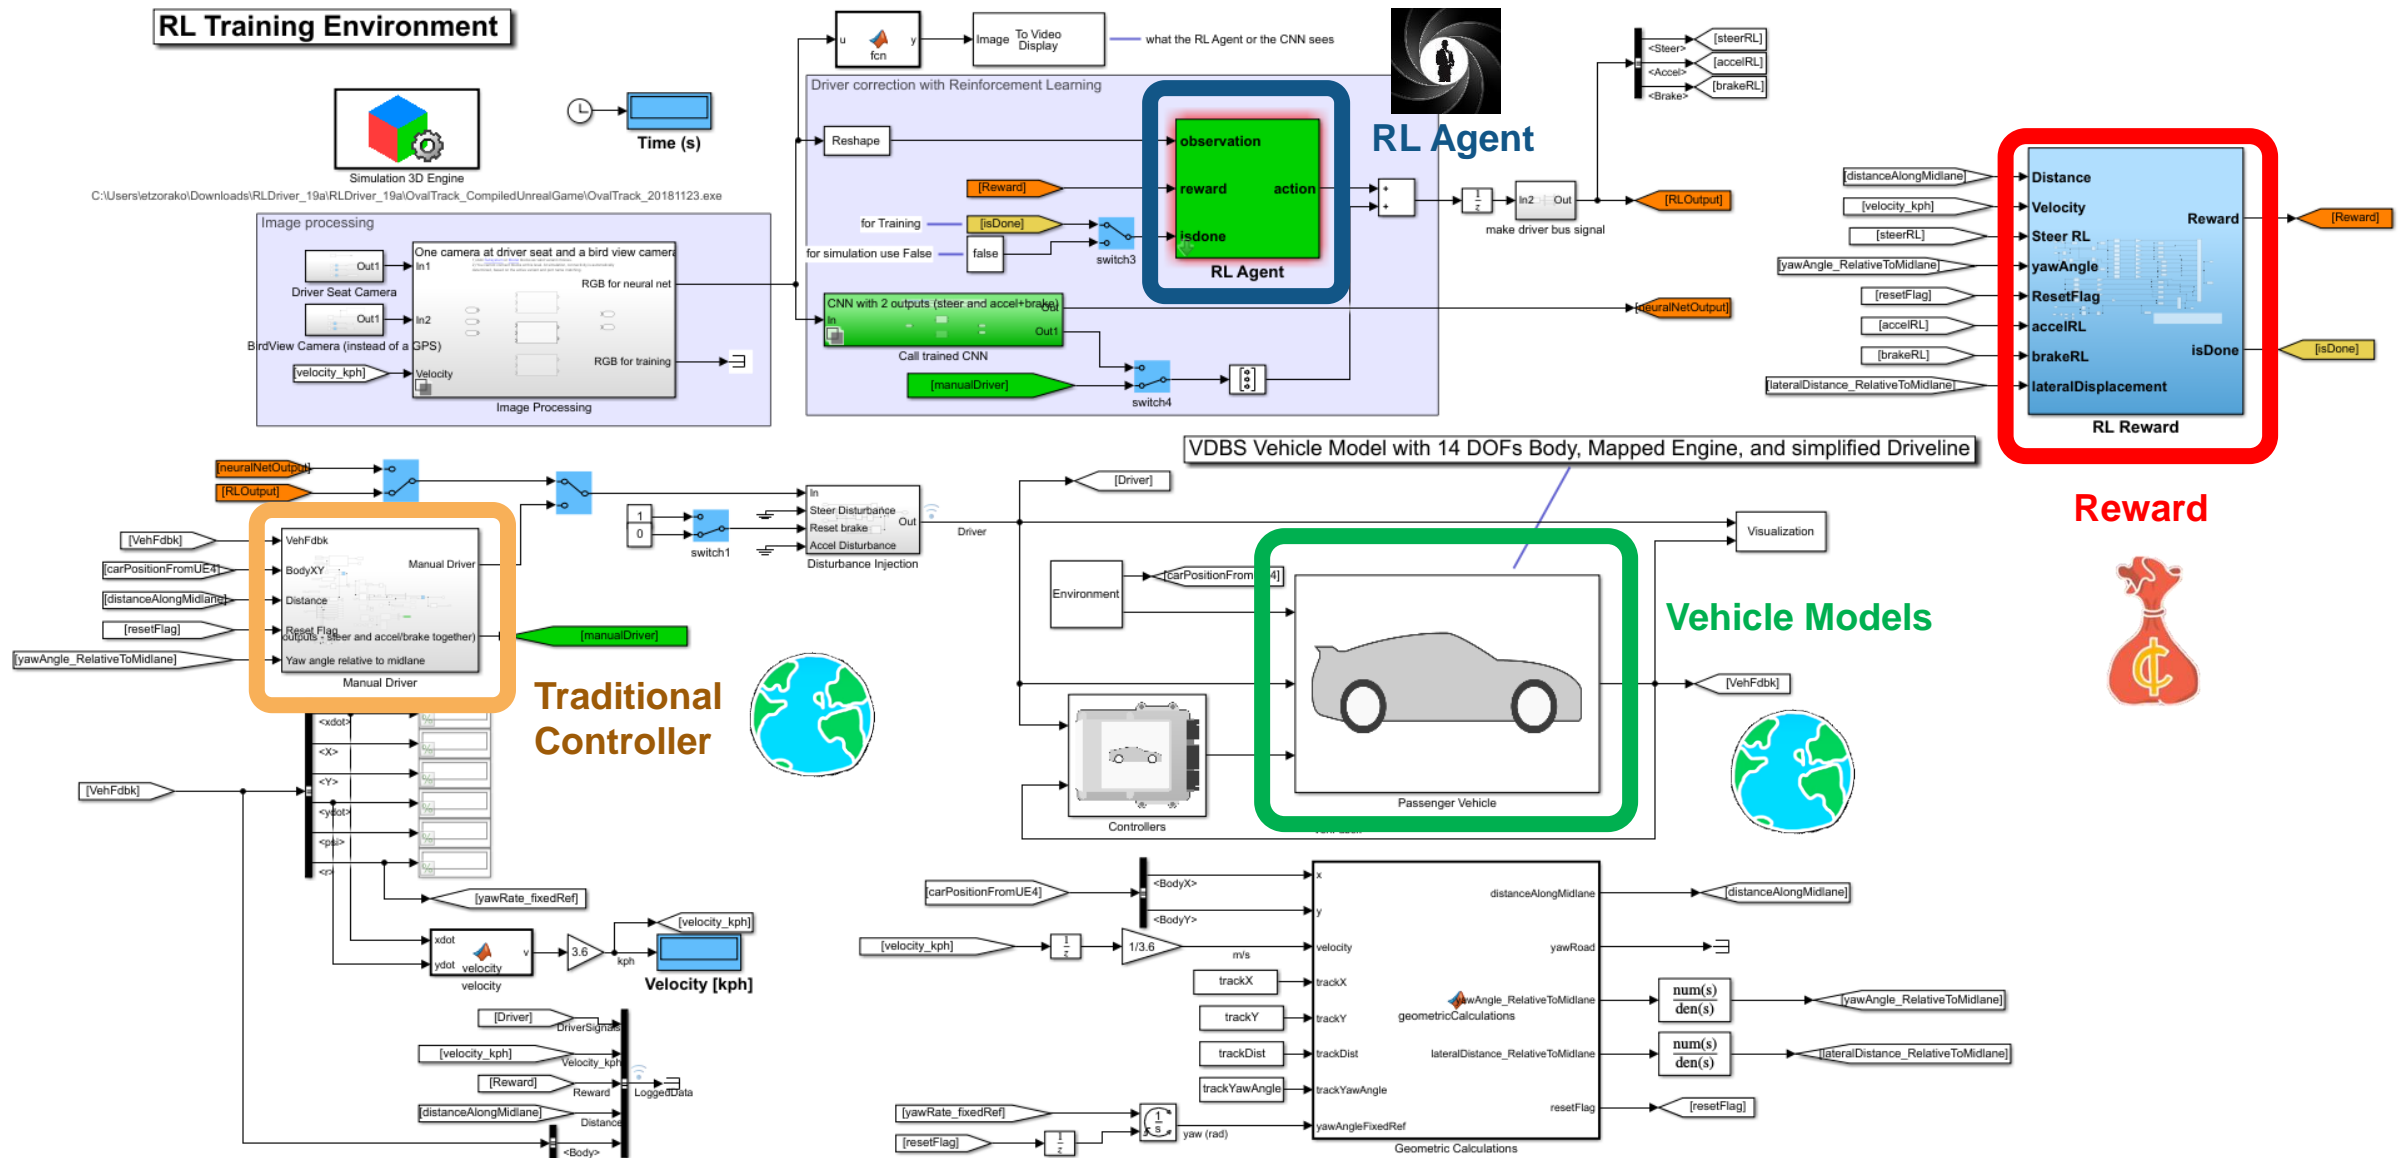
Apply Reinforcement Learning to **Control System Design** & **Decision Making Systems**
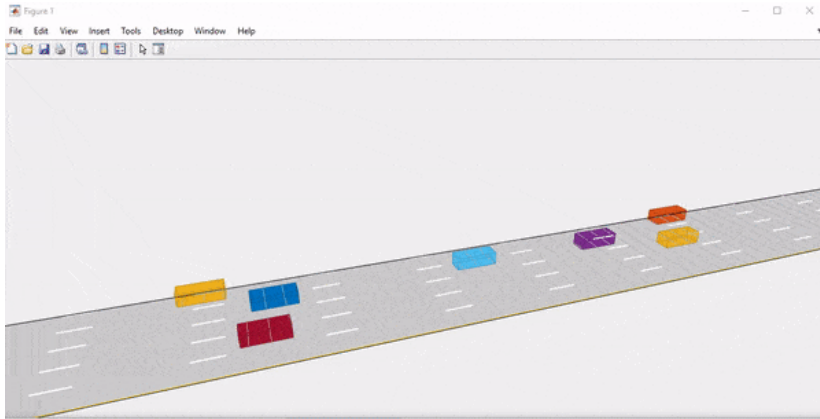
STATE-TABLES WITH VALUES

At each cell, four actions are possible: north, south, east, and west, which deterministically cause the agent to move one cell in the respective direction on the grid. Actions that would take the agent off the grid leave its location unchanged, but also result in a reward of −1. Other actions result in a reward of 0, except those that move the agent out of the special states A and B. From state A,

Gridworld

22.0	24.4	22.0	19.4	17.5
19.8	22.0	19.8	17.8	16.0
17.8	19.8	17.8	16.0	14.4
16.0	17.8	16.0	14.4	13.0
14.4	16.0	14.4	13.0	11.7

v_*

π_*

Policy for actions

From Sutton RL
1998

UMEÅ UNIVERSITY

Dynamic programming

- *Goal decomposition* as it solves complex high-level problems by decomposing them into smaller and smaller subproblems until it gets to a simple subproblem that can be solved

- Train a robot vacuum to move from one room in a house to its dock
    - stay in/exit this room?
    - move toward the door" or "move away from the door

- You need to know your house extremely well

DeepReinforcementLearningInAction-master/Chapter 1/Ch1_Introduction.ipynb

The Bellman equation for v

$$v_\pi(s) \doteq \mathbb{E}_\pi[G_t \mid S_t = s]$$

$$= \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} \mid S_t = s]$$

$$= \sum_a \pi(a|s) \sum_{s'} \sum_r p(s',r \mid s,a) \left[r + \gamma \mathbb{E}_\pi[G_{t+1}|S_{t+1} = s'] \right]$$

$$= \sum_a \pi(a|s) \sum_{s',r} p(s',r \mid s,a) \left[r + \gamma v_\pi(s') \right], \quad \text{for all } s \in \mathcal{S},$$

Monte Carlo

- The trial and error strategy
- Go to a party at a house that you've never been to before, you might have to look around until you find the bathroom on your own
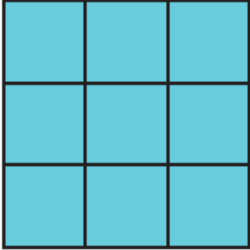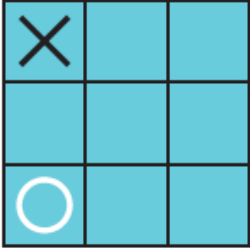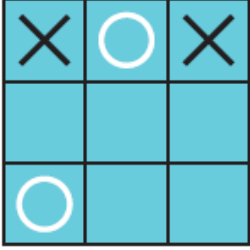
Agent

An action lookup table for Tic-Tac-Toe

there are 255,168 valid board positions..

Game play lookup table

Key Current state	Value Action to take
	Place × in top left
	Place × in top right
	Place × in bottom right

...

The multi-arm bandit problem

- 10 slot machines with a sign that says "Play for free! Max payout $10!

- Each have a different average payout
  - the reward at each play is probabilistic
  - so try to figure out which one gives the most rewards on average

- " play a few times, choosing different levers(a) and observing our rewards R for each action. Then we want to only choose the lever with the largest observed average reward."

$$Q_k(a) = \frac{R_1 + R_2 + ... + R_k}{k_a}$$

```
def exp_reward(action, history):
    rewards_for_action = history[action]
    return sum(rewards_for_action) /
        len(rewards_for_action)
```

action-value function; the value of taking a particular action

Math	Pseudocode

$\forall a_i \in A_k$	`def get_best_action(actions, history):`
	    `exp_rewards = [exp_reward(action, history) for action in`
	        `actions]`

$a* = \text{argmax}_a Q_k(a_i)$	`return argmax(exp_rewards)`

The following listing shows it as legitimate Python 3 code.

Listing 2.1 Finding the best actions given the expected rewards in Python 3

```python
def get_best_action(actions):
    best_action = 0
    max_action_value = 0
    for i in range(len(actions)):                         # Loops through all possible actions
        cur_action_value = get_action_value(actions[i])   # Gets the value of the current action
        if cur_action_value > max_action_value:
            best_action = i
            max_action_value = cur_action_value
    return best_action
```

a *greedy* (or exploitation) method!

Dilemma/Balance

Exploration

- Play the game and observe the rewards we get for the various machines
- To learn more
- Random selected action?

Exploitation

- Use our current knowledge about which machine seems to produce the most rewards, and keep playing that machine
- Greedy

Epsilon-greedy strategy

- with a probability, ε, we will choose an action, *a*, at random, and the rest of the time (probability 1 − ε) we will choose the best lever based on what we currently know from past plays

```
eps = 0.2
rewards = [0]
for i in range(500):
    if random.random() > eps:
        choice = get_best_arm(record)

    else:
        choice = np.random.randint(10)
    r = get_reward(probs[choice])
```

Chooses the best action with 0.8 probability, or randomly otherwise

Computes the reward for choosing the arm

Softmax selection policy

- Treating patients with heart attacks, choose 1 treatment out of 10
  - doesn't know which one is the best yet
  - randomly choosing a treatment could result in patient death, not just losing some money. We really want to make sure we don't choose the worst treatment

$$\Pr(A) = \frac{e^{Q_k(A)/\tau}}{\sum_{i=1}^{n} e^{Q_k(i)/\tau}}$$

```
def softmax(av, tau=1.12):
    softm = np.exp(av / tau) / np.sum( np.exp(av / tau) )
    return softm
```

τ is a parameter called *temperature* that scales the probability distribution of actions. Usually reduced over time, "cooling metal".

OPTIMISTIC INITIAL VALUE

Figure 2.3: The effect of optimistic initial action-value estimates on the 10-armed testbed. Both methods used a constant step-size parameter, $\alpha = 0.1$.

The Markov property

- the current state alone contains enough information to choose optimal actions to maximize future rewards

- A game (or any other control task) that exhibits the Markov property is said to be a *Markov decision process* (MDP):

Markov property or not?

- Driving a car

- Deciding whether to invest in a stock or not

- Choosing a medical treatment for a patient

- Diagnosing a patient's illness

- Predicting which team will win in a football game

- Choosing the shortest route (by distance) to some destination

- Aiming a gun to shoot a distant target

Episodic or Continuous task?

- most games are *episodic*, meaning that there are multiple chances to take actions before the game is over, and many games like chess don't naturally assign points to anything other than winning or losing the game

- a continuously learning RL algorithm has the advantage of being able to adapt to changing market conditions in real time
  - Need discount to avoid *infinite* reward

$$G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1},$$

maximize the expected discounted return Gt

SARSA

- $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha[r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$

- On-policy algorithm

TD-GAMMON (TESAURO 1992)

white pieces move counterclockwise

black pieces move clockwise

predicted probability of winning, $\hat{v}(S_t, \mathbf{w})$

TD error
$\hat{v}(S_{t+1}, \mathbf{w}) - \hat{v}(S_t, \mathbf{w})$

hidden units (40-80)

backgammon position (198 input units)
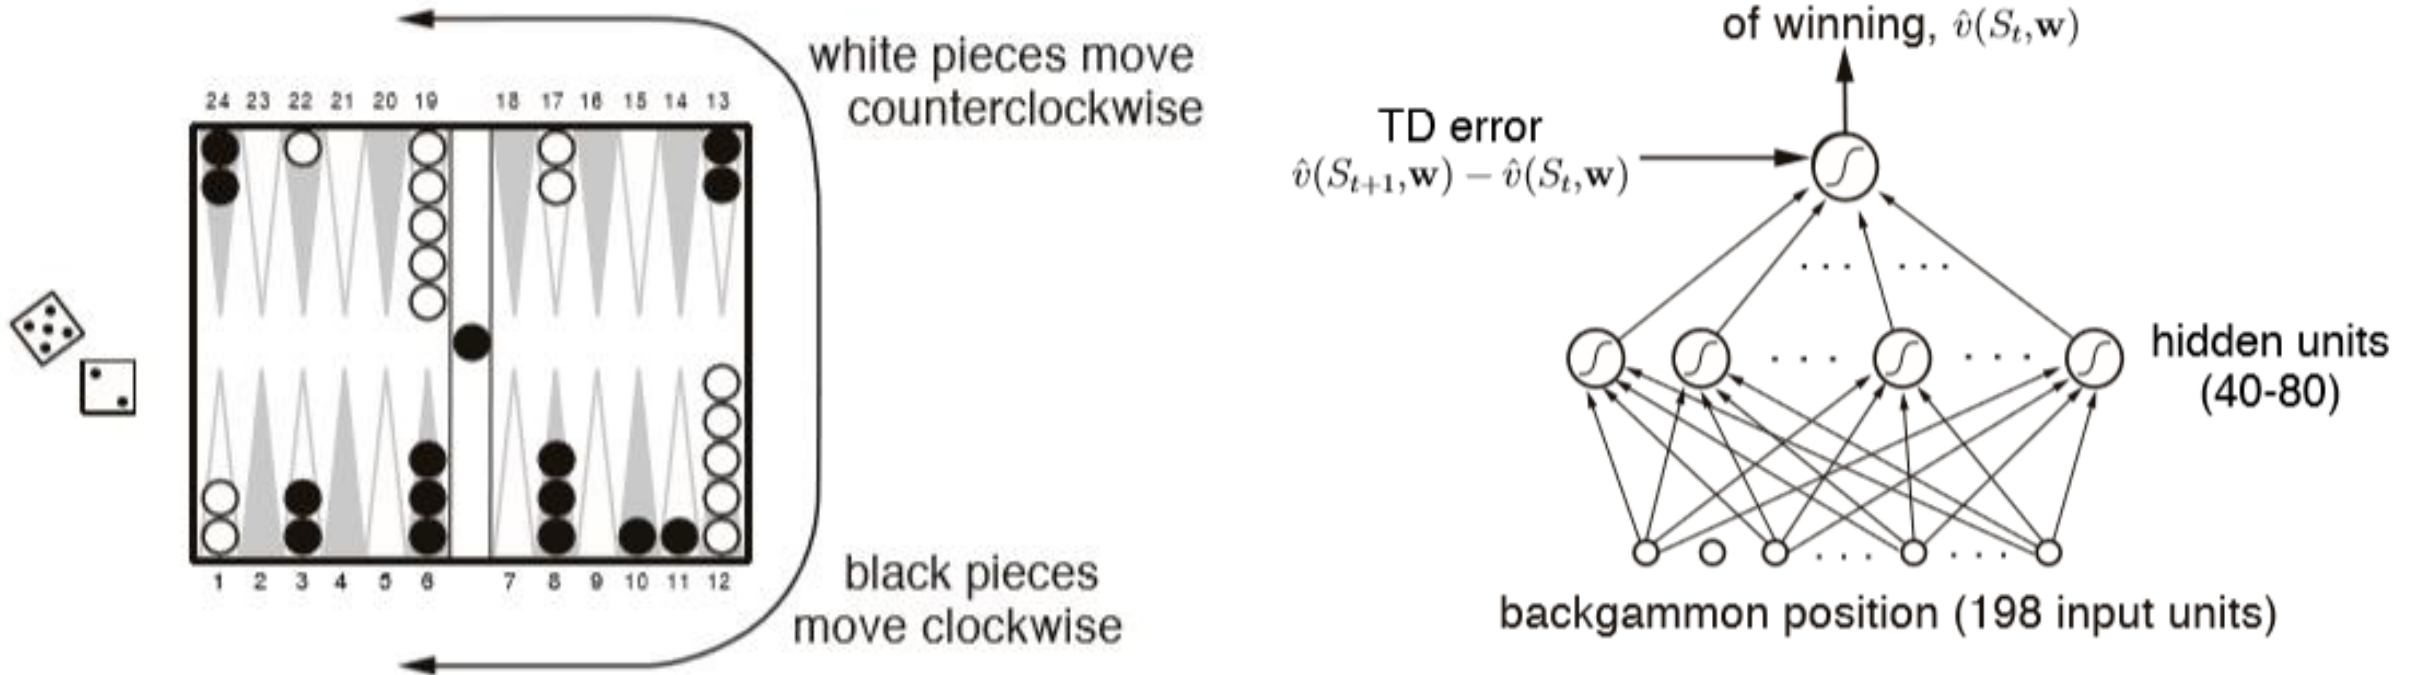
Figure 16.1: The TD-Gammon neural network

Q-learning algorithm update rule

Math

| Updated Q value | Current Q value | Observed reward | Max Q value for all actions | = Off-policy algorithm |

$$Q(S_t, A_t) = Q(S_t, A_t) + \alpha[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

Step size   Discount factor

Pseudocode

```
def get_updated_q_value(old_q_value, reward, state, step_size, discount):
    term2 = (reward + discount * max([Q(state, action) for action in
            actions]))
    term2 = term2 - old_q_value
    term2 = step_size * term2
    return (old_q_value + term2)
```

English

The Q value at time *t* is updated to be the current predicted Q value plus the amount of value we expect in the future, given that we play optimally from our current state.

On-policy algorithm

Off-policy algorithm

ANN   DEEP LEARNING

Patterns of Local Contrast

Face Features

Face

Input Layer

Hidden Layer 1

Hidden Layer 2

Output Layer

RL ♥ DL

- Combine RL and DL
- Replace discrete tables with interpolation function
  - Can generalize better, don't need to visit every state
  - Learns faster/better
  - Problems: Deadly Triad, Catastrophic forgetting, Complexity..
- Why now?
  - Powerful computers
  - Multicore, GPU board (Nvidia Geforce RTX 2080 Ti)
  - Libraries; TensorFlow, Matlab, PyTorch

UMEÅ UNIVERSITY

Contextual bandits

- Example: advertisement placement
  - maximize the probability that you will click ads depending on what site you visit
  - Let's say we manage 10 e-commerce websites, each focusing on selling a different broad category of retail items such as computers, shoes, jewelry, etc.

We want to increase sales by referring customers who shop on one of our sites to another site that they might be interested in. When a customer checks out on a particular site in our network, we will display an advertisement to one of our other sites in hopes they'll go there and buy something else.

Alternatively, we could place an ad for another product on the same site. Our problem is that we don't know which sites we should refer users to. We could try placing random ads, but we suspect a more targeted approach is possible.

state spaces

- we know the user is buying something on a particular site, which may give us some information about that user's preferences and could help guide our decision about which ad to place
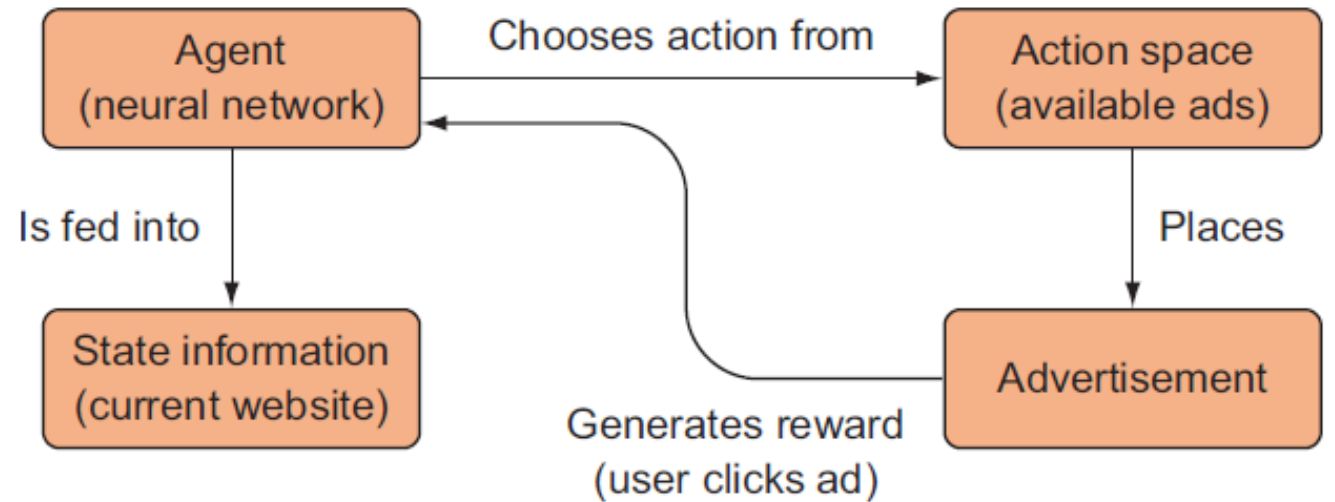
Building networks with PyTorch

A simple two-layer neural network with an

Optimizer:

```
model = torch.nn.Sequential(

    torch.nn.Linear(10, 150),

    torch.nn.ReLU(),

    torch.nn.Linear(150, 4),

    torch.nn.ReLU(),

)
```

A training loop:
```
for step in range(100):
    y_pred = model(x)
    loss = loss_fn(y_pred, y_correct)
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
```

Our own Python class:
```
from torch.nn import Module, Linear
class MyNet(Module):
    def __init__(self):
        super(MyNet, self).__init__()
        self.fc1 = Linear(784, 50)
        self.fc2 = Linear(50, 10)
    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
    return x
model = MyNet()
```

Exercice 1

- https://colab.research.google.com/
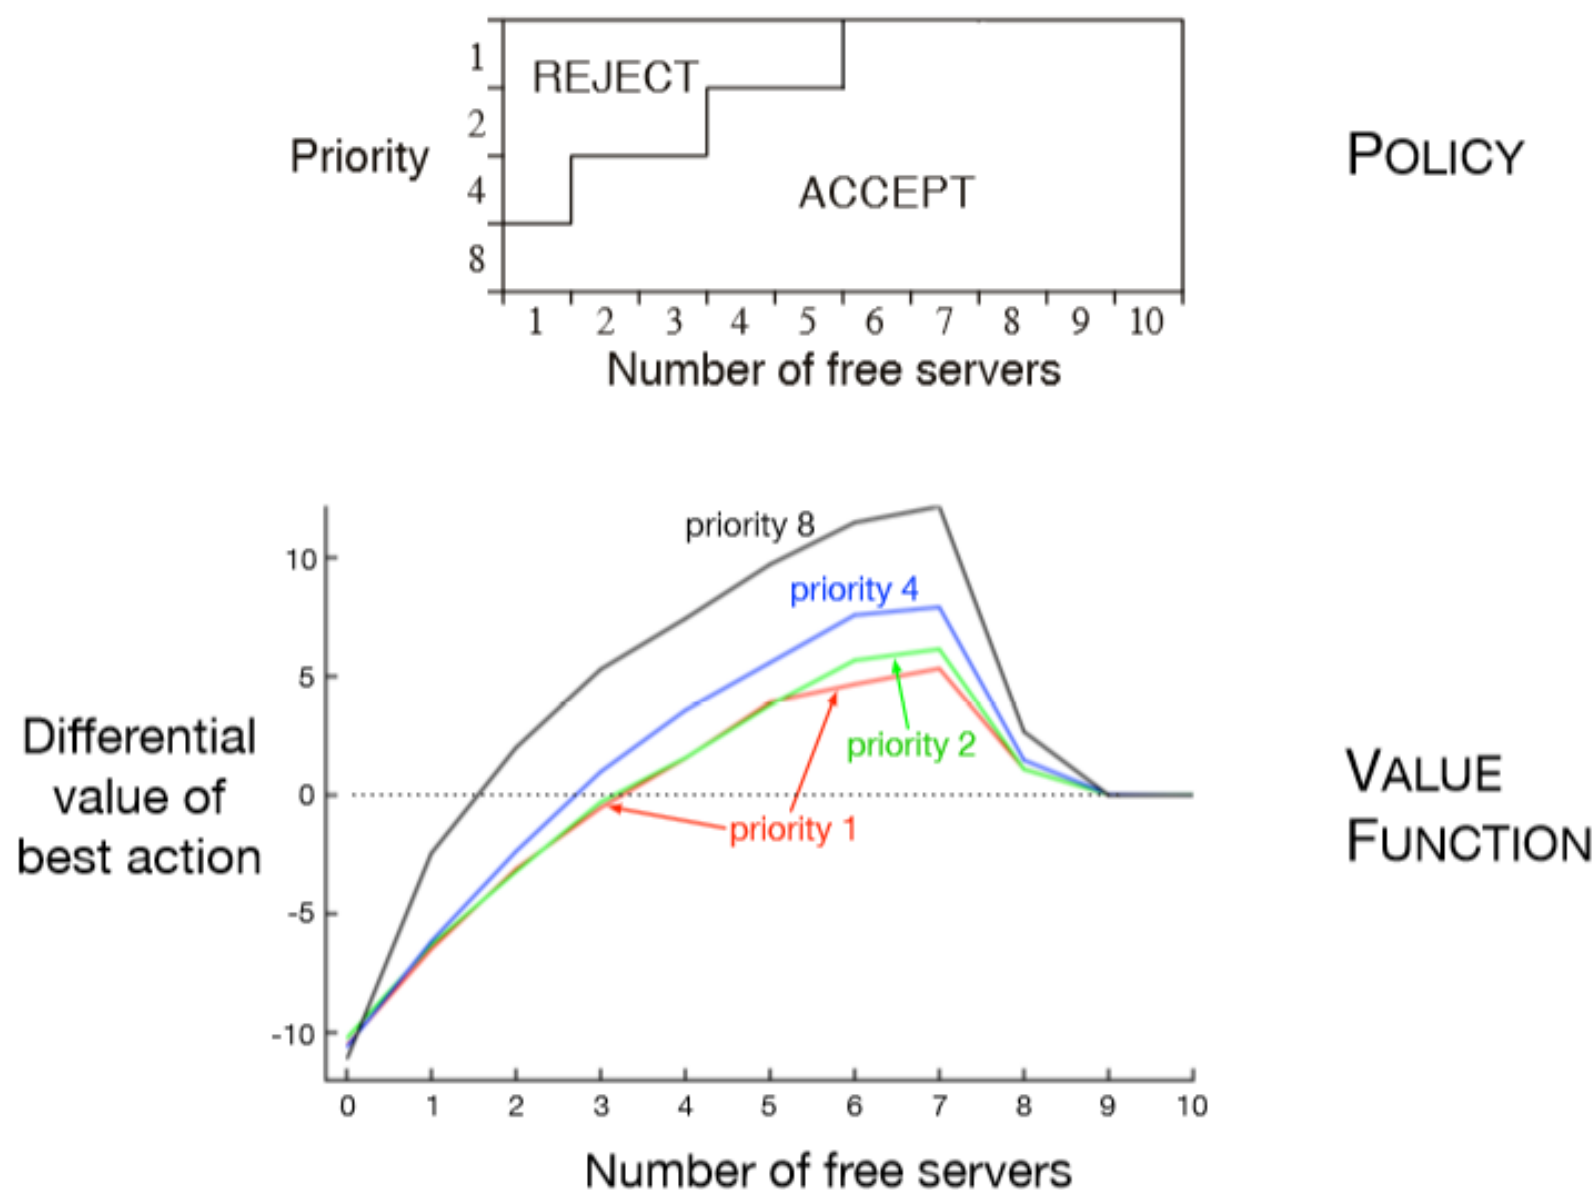- https://github.com/nutte2/DeepRL
- bandits

Figure 10.5: The policy and value function found by differential semi-gradient one-step Sarsa on the access-control queuing task after 2 million steps. The drop on the right of the graph is probably due to insufficient data; many of these states were never experienced. The value learned for $\bar{R}$ was about 2.31. ∎

Example 5.3: Solving Blackjack   It is straightforward to apply Monte Carlo ES to blackjack. Because the episodes are all simulated games, it is easy to arrange for exploring starts that include all possibilities. In this case one simply picks the dealer's cards, the player's sum, and whether or not the player has a usable ace, all at random with equal probability. As the initial policy we use the policy evaluated in the previous blackjack example, that which sticks only on 20 or 21. The initial action-value function can be zero for all state–action pairs. Figure 5.2 shows the optimal policy for blackjack found by Monte Carlo ES. This policy is the same as the "basic" strategy of Thorp (1966) with the sole exception of the leftmost notch in the policy for a usable ace, which is not present in Thorp's strategy. We are uncertain of the reason for this discrepancy, but confident that what is shown here is indeed the optimal policy for the version of blackjack we have described.

Figure 5.2: The optimal policy and state-value function for blackjack, found by Monte Carlo ES (Figure 5.4). The state-value function shown was computed from the action-value function found by Monte Carlo ES.

Figure 10.1: The Mountain Car task (upper left panel) and the cost-to-go function $(-\max_a \hat{q}(s, a, \mathbf{w}))$ learned during one run.

Open AI Gym

```
import gym
env = gym.make('CarRacing-v0')
env.reset()
env.step(action)
env.render()
```

Exercice 2

- https://colab.research.google.com/
- https://github.com/nutte2/DeepRL
- mountcar

GOOGLE DEEPMINDS ALPHAGO ZERO

- Beat the previous version of AlphaGo (Final score: 100–0).
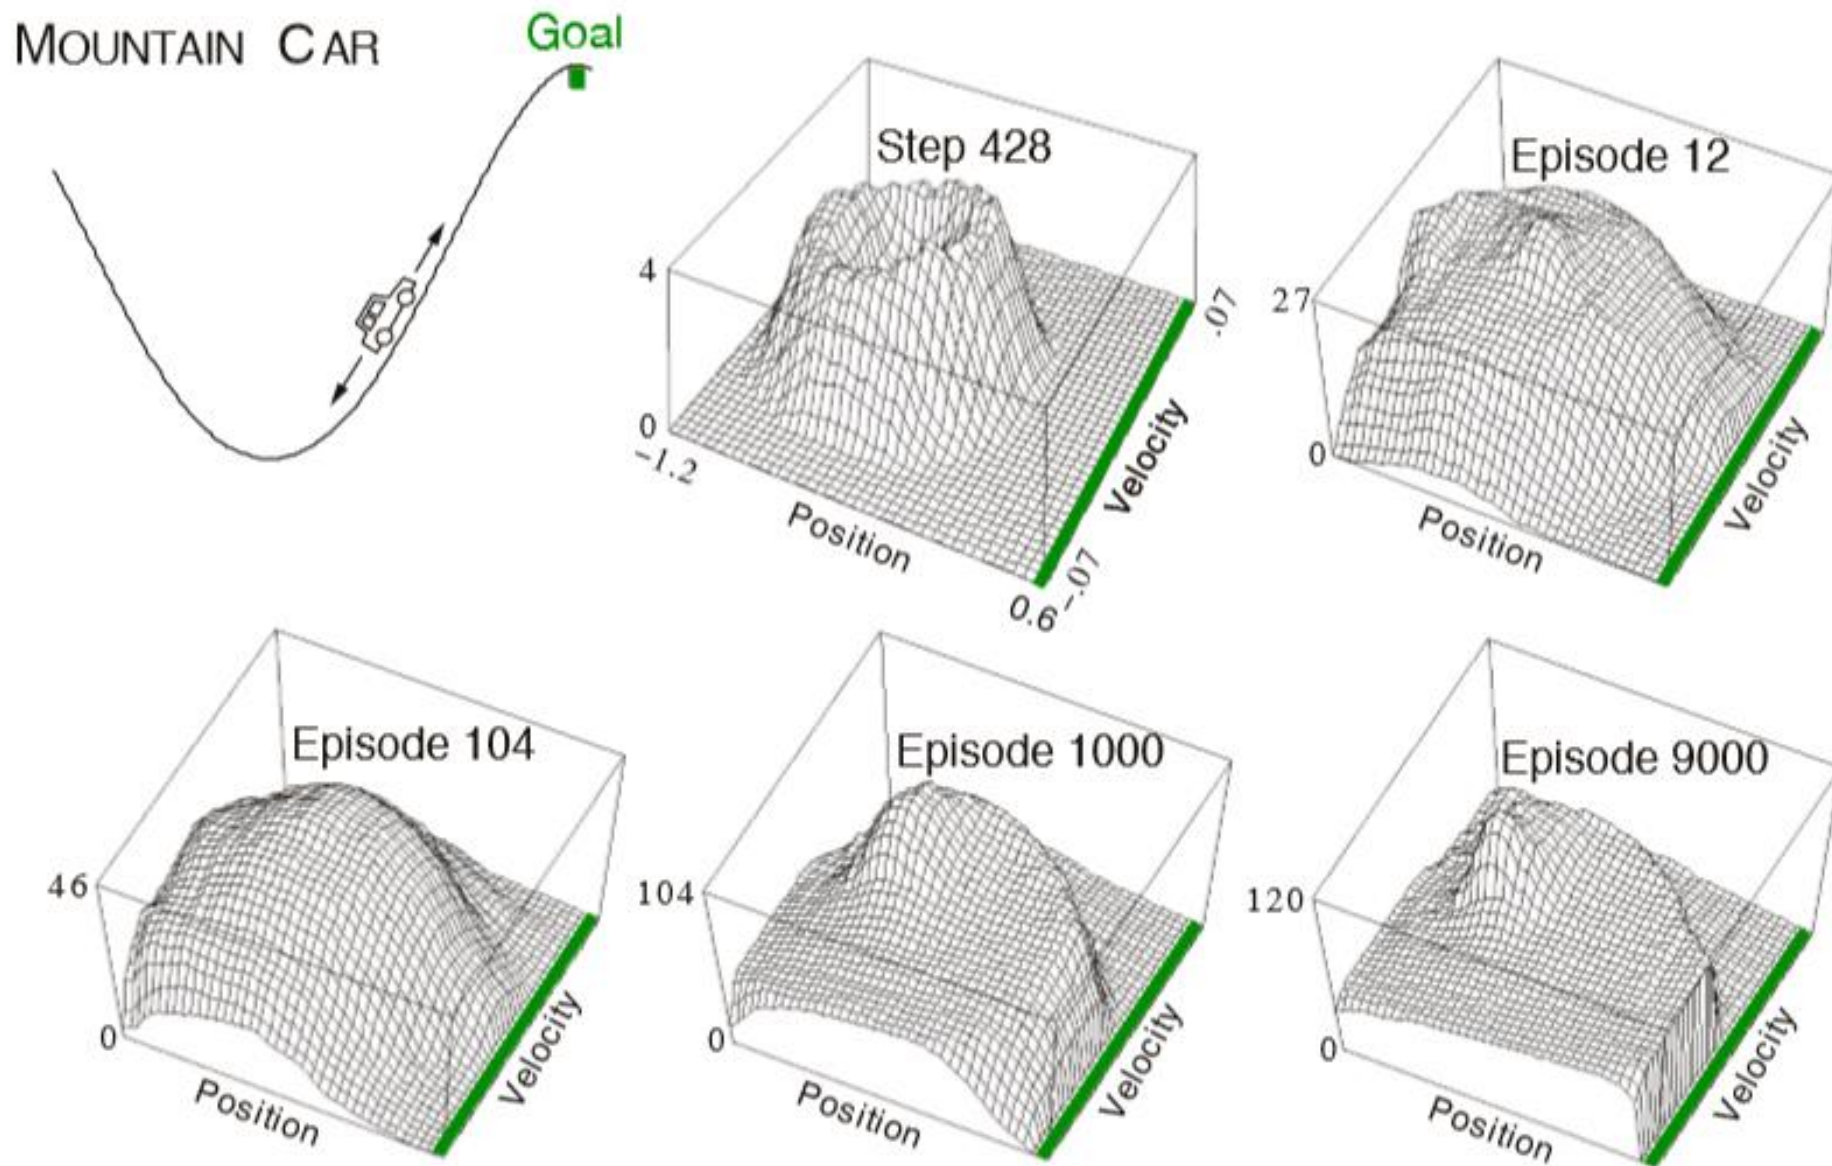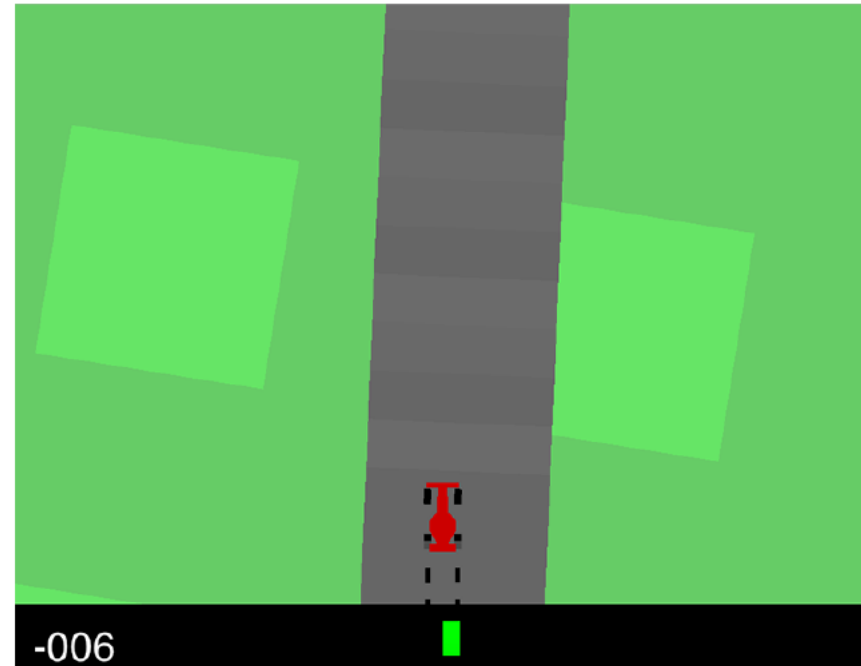- Learn to perform this task from scratch, without learning from previous human knowledge (i.e. recorded game play).
- World champion level Go playing in just 3 days of training.
- Do so with an order of magnitude less neural networks (4 TPUs vs 48 TPUs).
- Do this with less training data (3.9 million games vs 30 millions games).
- https://deepmind.com/blog/article/alphago-zero-starting-scratch

UMEÅ UNIVERSITY

DeepMind's DQN algorithm for Atari games

Four (three here) 84 × 84 grayscale images at each step, which would lead to 256^{28228} unique game states. DQN CNN have 1792 parameters

Policy functions

- How exactly do we use our current state information to decide what action to take?

- In words, a policy, π, is the strategy of an agent in some environment.

- A policy is a function that maps a state to a probability distribution over the set of possible actions in that state

Math	English
$\pi, s \rightarrow Pr(A \mid s)$, where $s \in S$	A policy, π, is a mapping from states to the (probabilistically) best actions for those states.

- The *optimal policy*—it's the strategy that maximizes rewards.

Math	English
$\pi* = argmax\ E(R \mid \pi)$,	If we know the expected rewards for following any possible policy, π, the optimal policy, $\pi*$, is a policy that, when followed, produces the maximum possible rewards.

Value functions

- *Value functions* are functions that map a state to the *expected value* (the expected reward) of being in some state

Math	English
$V_\pi: s \rightarrow E(R \mid s, \pi),$	A value function, V_π, is a function that maps a state, s, to the expected rewards, given that we start in state s and follow some policy, π.

- Accepts a state, *s*, and returns the expected reward of starting in that state and taking actions according to our policy, π
- Actual values depends on the policy

Q function or Q value

- Estimates of the expected reward for taking an action given a state

| Math	English
$Q_\pi: (s\|a) \rightarrow E(R\|a,s,\pi),$	Q_π is a function that maps a pair, (s, a), of a state, s, and an action, a, to the expected reward of taking action a in state s, given that we're using the policy (or "strategy") π.

- Q-learning (Watkins 1989, 1992)
- https://en.wikipedia.org/wiki/Q-learning

Agent – Action – Environment – State/Reward

A simple version of Gridworld

- Train a DRL agent to navigate the Gridworld board to the goal
- Following the most efficient route

A : Agent
W : Wall
— : Pit
+ : Goal

Sequence of events for a game of Gridworld

- We start the game in some state that we'll call S_t. The state includes all the information about the game that we have. For our Gridworld example, the game state is represented as a 4 × 4 × 4 tensor.

- We feed the S_t data and a candidate action into a deep neural net and it produces a prediction of how valuable taking that action in that state is

- The algorithm is not predicting the reward we will get after taking a particular action; it's predicting the expected value (the expected rewards), which is the long-term average reward we will get from taking an action in a state and then continuing to behave according to our policy π.

We take an action,

- perhaps because our neural network predicted it is the highest value action or perhaps we take a random action. We'll label the action A_t. We are now in a new state of the game, which we'll call S_{t+1}, and we receive or observe a reward, labelled R_{t+1}. We want to update our learning algorithm to reflect the actual reward we received, after taking the action it predicted was the best. Perhaps we got a negative reward or a really big reward, and we want to improve the accuracy of the algorithm's predictions

Updating Q(S_t, a)

- Now we run the algorithm using S_{t+1} as input and figure out which action our algorithm predicts has the highest value. We'll call this value max $Q(S_{t+1}, a)$. To be clear, this is a single value that reflects the highest predicted Q value, given our new state and all possible actions there

- We'll perform one iteration of training using some loss function, such as mean squared error, to minimize the difference between the predicted value from our algorithm and the target prediction of

$$Q(S_t, A_t) + \alpha *[R_{t+1} + \gamma *\max_a Q(S_{t+1}, A) - Q(S_t, A_t)]$$

DeepMind used a modified vector-valued Q function that accepts a state and returns a vector of state-action values, one for each possible action given the input state. The vector-valued Q function is more efficient, since you only need to compute the function once for all the actions
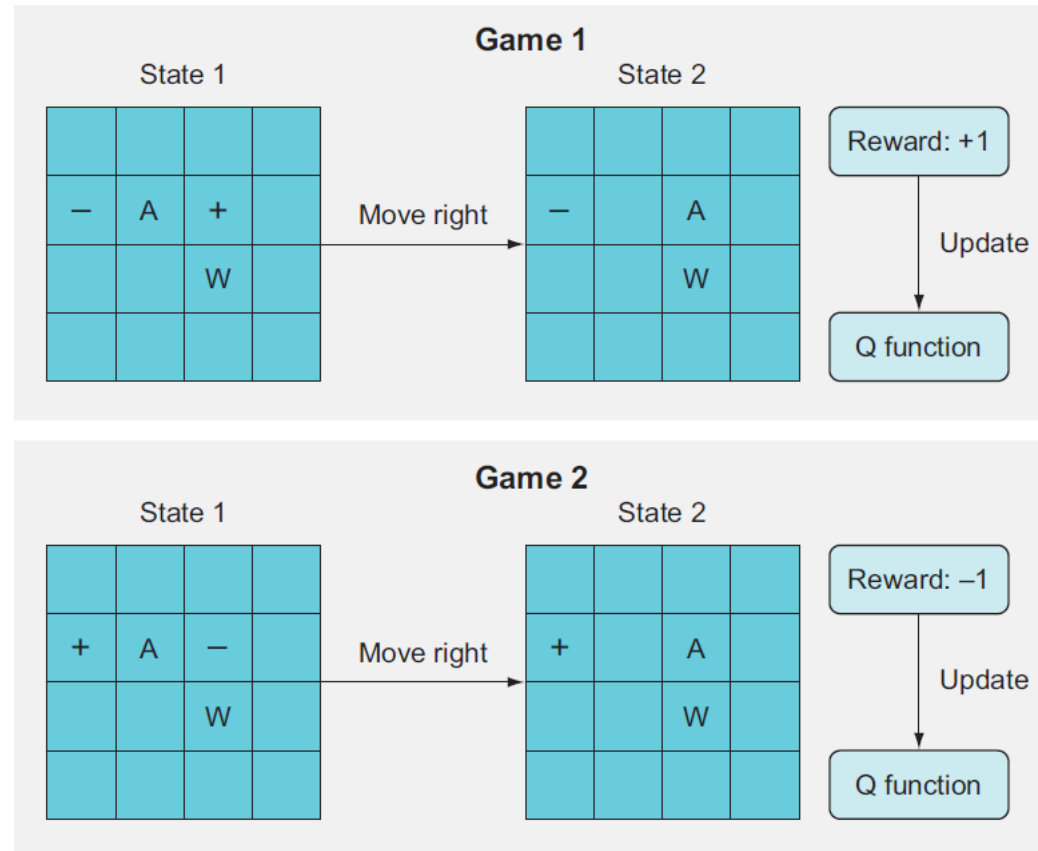
The Gridworld board is represented as a numpy array. It is a 4 x 4 x 4 tensor, composed of 4 "slices"

The general architecture for the model we will build.

Catastrophic forgetting

two game states are very similar and yet lead to very different outcomes, the Q function will get "confused" and won't be able to learn what to do

Preventing catastrophic forgetting: Experience replay

- employ mini-batching by storing past experiences and then using a random subset of these experiences to update the Q-network, rather than using just the single most recent experience

Experience replay
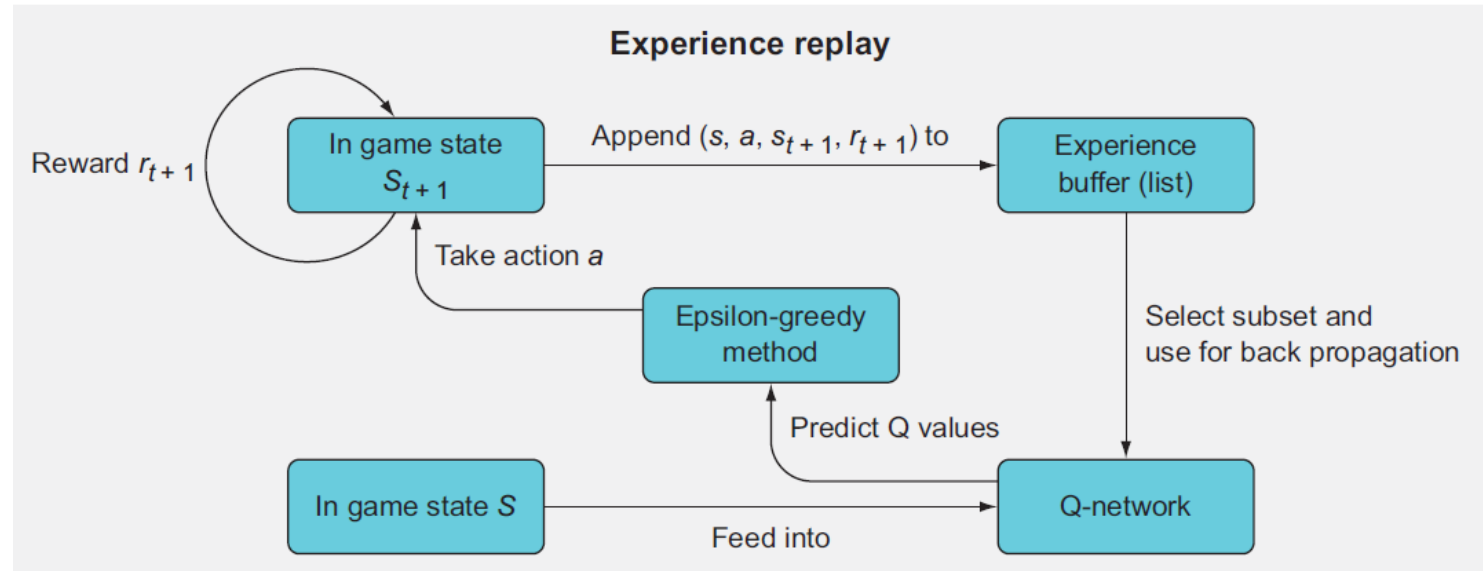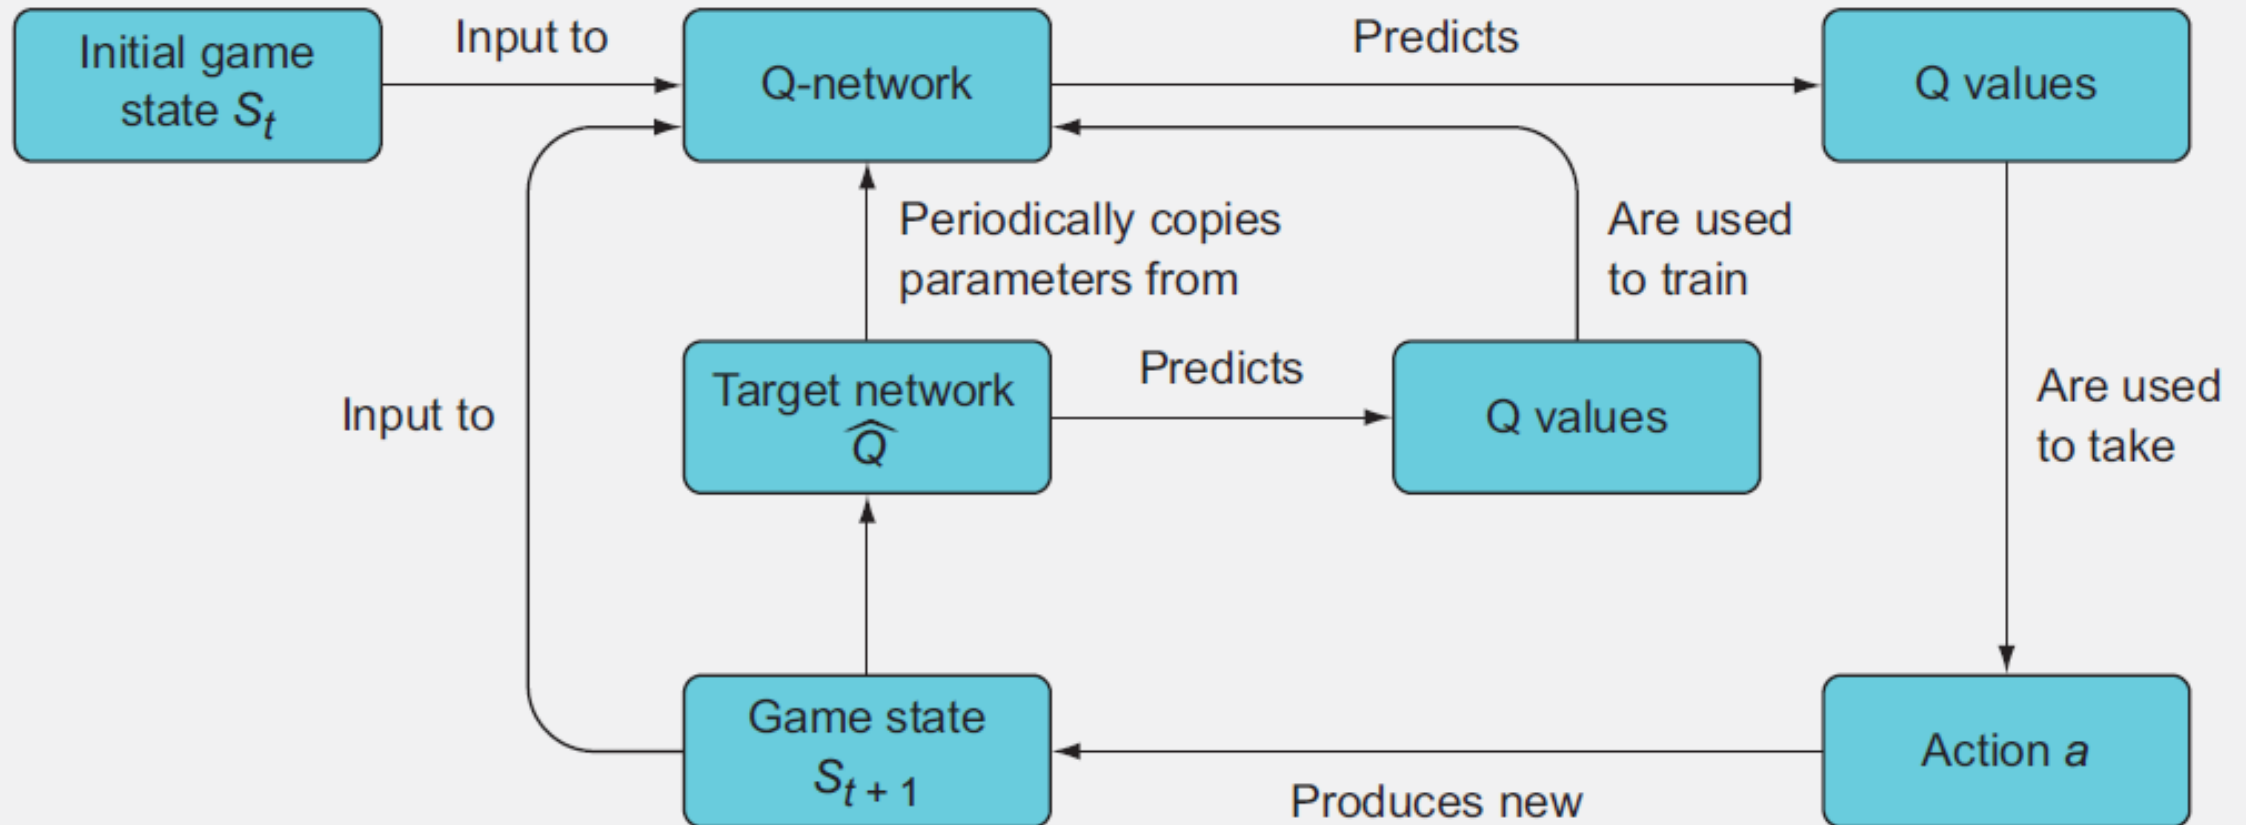
Improving stability with a target network

Exercice 3

- https://colab.research.google.com/
- https://github.com/nutte2/DeepRL
- grids

Some DRL links

- Deep Reinforcement Learning: Pong from Pixels
    - http://karpathy.github.io/2016/05/31/rl/
    - https://arxiv.org/pdf/1312.5602.pdf
- the DRL book's GitHub repository: http://mng.bz/JzKp
- Tensorflow TF-Agents
    - https://colab.research.google.com/github/tensorflow/agents/blob/master/docs/tutorials/1_dqn_tutorial.ipynb
- Real application: Adaptive Power System Emergency Control using Deep Reinforcement Learning
    - https://arxiv.org/pdf/1903.03712.pdf